

Replicating The Log of Gravity

Mauricio “Pachá” Vargas Sepúlveda
Department of Political Science, University of Toronto

Last updated: 2024-09-05 00:30

Contents

1	About	1
2	Obtaining the original codes and data	1
3	R packages	1
4	Loading the original data	1
5	Poisson Pseudo Maximum Likelihood	2
6	Ordinary Least Squares	2
7	Tobit	3
8	Non-Linear Least Squares	4
9	Putting it all together	4
10	Acknowledgements	6
	References	6

1 About

Here I replicate the main results from Santos Silva and Tenreyro (2006) in R. The original results were obtained in TSP back in 2006.

The idea here is to be explicit regarding the conceptual approach to regression in R. For most of the replication I used base R (R Core Team 2021) without external libraries (i.e packages) except when it was absolutely necessary.

Much of the methods exposed here lead to the exact same results as using the gravity package (Woelwer et al. 2020) which provides convenient wrappers for gravity estimation.

2 Obtaining the original codes and data

I shall organize the original codes and data from the authors' site to put these on GitHub and therefore ease reproducibility in case of broken links or anything that makes it difficult to obtain the original zip file with the data and codes.

```
url <- "https://personal.lse.ac.uk/tenreyro/regressors.zip"
zip <- gsub(".*/", "", url)
if (!file.exists(zip)) try(download.file(url, zip))

dout <- "regressors"
if (!dir.exists(dout)) unzip(zip, exdir = dout)
```

3 R packages

```
library(haven) # read Stata datasets
library(censReg) # Tobit estimation
library(stargazer) # table of results
```

4 Loading the original data

Thanks to the haven package (Wickham and Miller 2021) we can read Stata datasets directly in R without loss of information about column types and other common problems when reading proprietary formats.

```
log_of_gravity <- read_dta(paste0(dout, "/Log of Gravity.dta"))
```

5 Poisson Pseudo Maximum Likelihood

Table 3 in Santos Silva and Tenreyro (2006) summarises a large portion of the article. In R we would replicate it by fitting two Generalized Linear Models since the article introduces estimates with and without removing zero flows.

```
ppml_formula <- trade ~ lypex + lypim + lyex + lyim + ldist + border +  
  comlang + colony + landl_ex + landl_im + lremot_ex + lremot_im +  
  comfrt_wto + open_wto  
  
fit_ppml_1 <- glm(  
  ppml_formula,  
  data = log_of_gravity,  
  subset = trade > 0,  
  family = quasipoisson()  
)  
  
fit_ppml_2 <- glm(  
  ppml_formula,  
  data = log_of_gravity,  
  family = quasipoisson()  
)
```

The replication effort here is null, it just sufficed to look at the summary table in the article and subset the data to drop zero flows. Therefore, it makes sense to proceed with the other models.

6 Ordinary Least Squares

The only consideration here is to drop zero flows for some of the models with log in the dependent variable even when Table 3 is not explicit about this, otherwise we break the fitting algorithm.

For example, for estimations of the type $\log(\text{trade}) = \beta_0 + \beta_1 \text{lypex} + \dots + \varepsilon$, we need to drop zero flows to replicate the result. On the other hand, for estimations of the type $\log(1 + \text{trade}) = \beta_0 + \beta_1 \text{lypex} + \dots + \varepsilon$, we don't need to drop zero flows.

```
fit_ols_1 <- lm(  
  update.formula(ppml_formula, log(.) ~ .),  
  data = log_of_gravity,  
  subset = trade > 0  
)  
  
fit_ols_2 <- lm(  
  update.formula(ppml_formula, log(1 + .) ~ .),  
  data = log_of_gravity
```

)

7 Tobit

The Tobit estimation is similar but requires the use of the `censReg` package (Henningsen 2020). The complicated part of the estimation is to extract the right hand side of the model formula to define a vector of zeroes with a length equal to the right hand side plus two as starting point for the Maximum Likelihood estimation (e.g., including the depending variable and intercept besides the estimating slopes).

In order to obtain the a value that matches the results in the article I proceeded with an iteration loop until achieving convergence with respect to one of the estimated slopes. The initial value of $a = 200$ was arbitrary and set after trying reasonable guesses that converge to the slopes in the original article after 9 iterations for a final value of $a = 159$.

```
a <- 200
lypex_ref <- 1.058
tol <- 0.001
lypex_estimate <- 2 * lypex_ref
iter <- 0

while (abs(lypex_estimate - lypex_ref) > tol) {
  log_of_gravity$log_trade_cens <- log(a + log_of_gravity$trade)
  log_trade_cens_min <- min(log_of_gravity$log_trade_cens, na.rm = TRUE)

  fit_tobit <- censReg(
    formula = update.formula(ppml_formula, log_trade_cens ~ .),
    left = log_trade_cens_min,
    right = Inf,
    data = log_of_gravity,
    start = rep(0, 2 + length(attr(terms(ppml_formula), "term.labels"))),
    method = "BHHH"
  )

  lypex_estimate <- coef(fit_tobit)[2]
  if (abs(lypex_estimate - lypex_ref) > 2 * tol) {
    a <- a - 5
  } else {
    a <- a - 1
  }
  iter <- iter + 1
}
```

8 Non-Linear Least Squares

For this type of estimation the starting values are retrieved from the results of the PPML model with zero flows and then we pass these values to a Generalized Linear Model using the Gaussian distribution and a log-link.

```
fit_ppml_eta <- fit_ppml_2$linear.predictors
fit_ppml_mu <- fit_ppml_2$fitted.values
fit_ppml_start <- fit_ppml_2$coefficients

fit_nls <- glm(
  ppml_formula,
  data = log_of_gravity,
  family = gaussian(link = "log"),
  etastart = fit_ppml_eta,
  mustart = fit_ppml_mu,
  start = fit_ppml_start,
  control = list(maxit = 200, trace = FALSE)
)
```

9 Putting it all together

There wasn't much effort involved in the replication, which is something desirable. I didn't even have to email the authors with questions whereas the data was filtered or transformed in ways not mentioned in the article, which is something that we often see. The article is very close to full replication according to the criteria from Peng (2011).

```
stargazer(
  fit_ols_1, fit_ols_2, fit_tobit, fit_nls, fit_ppml_1, fit_ppml_2,
  header = FALSE, font.size = "footnotesize", model.names = F,
  omit.table.layout = "d", omit.stat = c(
    "f", "ser", "ll", "aic", "bic", "rsq", "adj.rsq"
  ),
  title = "Replication results for OLS (1-2), Tobit (3), NLS (4) and
  PPML (5-6).",
)
```

Table 1: Replication results for OLS (1-2), Tobit (3), NLS (4) and PPML (5-6).

	<i>Dependent variable:</i>					
	(1)	(2)	(3)	(4)	(5)	(6)
lypex	0.938*** (0.012)	1.128*** (0.011)	1.059*** (0.011)	0.738*** (0.004)	0.721*** (0.008)	0.732*** (0.006)
lypim	0.798*** (0.011)	0.866*** (0.011)	0.848*** (0.010)	0.862*** (0.005)	0.732*** (0.008)	0.741*** (0.006)
lyex	0.207*** (0.017)	0.277*** (0.017)	0.228*** (0.014)	0.396*** (0.010)	0.154*** (0.013)	0.157*** (0.010)
lyim	0.106*** (0.017)	0.217*** (0.017)	0.178*** (0.014)	-0.033*** (0.007)	0.133*** (0.013)	0.135*** (0.010)
ldist	-1.166*** (0.034)	-1.151*** (0.037)	-1.160*** (0.029)	-0.924*** (0.008)	-0.776*** (0.018)	-0.784*** (0.013)
border	0.314** (0.143)	-0.241 (0.164)	-0.225** (0.109)	-0.081*** (0.010)	0.202*** (0.034)	0.193*** (0.026)
comlang	0.678*** (0.064)	0.742*** (0.064)	0.759*** (0.052)	0.689*** (0.016)	0.751*** (0.037)	0.746*** (0.028)
colony	0.397*** (0.068)	0.392*** (0.068)	0.416*** (0.056)	0.036** (0.018)	0.020 (0.043)	0.025 (0.032)
landl_ex	-0.062 (0.065)	0.106* (0.060)	-0.038 (0.060)	-1.367*** (0.031)	-0.872*** (0.057)	-0.863*** (0.043)
landl_im	-0.665*** (0.063)	-0.278*** (0.060)	-0.478*** (0.059)	-0.471*** (0.022)	-0.703*** (0.054)	-0.696*** (0.040)
lremot_ex	0.467*** (0.078)	0.526*** (0.089)	0.563*** (0.077)	1.188*** (0.018)	0.647*** (0.048)	0.660*** (0.036)
lremot_im	-0.205** (0.081)	-0.109 (0.089)	-0.032 (0.074)	1.010*** (0.018)	0.549*** (0.048)	0.562*** (0.036)
comfrt_wto	0.491*** (0.105)	1.289*** (0.143)	0.728*** (0.113)	0.443*** (0.014)	0.179*** (0.036)	0.181*** (0.027)
open_wto	-0.170*** (0.049)	0.739*** (0.048)	0.310*** (0.040)	0.928*** (0.024)	-0.139*** (0.039)	-0.107*** (0.029)
logSigma			0.677*** (0.007)			
Constant	-28.492*** (1.088)	-39.909*** (1.221)	-36.626*** (1.059)	-45.098*** (0.239)	-31.530*** (0.596)	-32.326*** (0.444)
Observations	9,613	18,360	18,360	18,360	9,613	18,360

Note:

10 Acknowledgements

Thanks to Joao Santos Silva for pointing that out, in a previous version of this document I mentioned that the original results were obtained with Stata.

References

- Henningsen, Arne. 2020. *censReg: Censored Regression (Tobit) Models*. <https://CRAN.R-project.org/package=censReg>.
- Peng, Roger D. 2011. “Reproducible Research in Computational Science.” *Science* 334 (6060): 1226–27.
- R Core Team. 2021. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org/>.
- Santos Silva, Joao, and Silvana Tenreiro. 2006. “The Log of Gravity.” *The Review of Economics and Statistics* 88 (4): 641–58.
- Wickham, Hadley, and Evan Miller. 2021. *haven: Import and Export SPSS, Stata and SAS Files*. <https://CRAN.R-project.org/package=haven>.
- Woelwer, Anna-Lena, Jan Pablo Burgard, Joshua Kunst, and Mauricio Vargas. 2020. *Gravity: Estimation Methods for Gravity Models*. <http://pacha.dev/gravity>.